

Introduction to Sockets

CS32: Spring 2010

1 Introduction

In this lab you are going to learn about socket programming in Java and write your own client/server application, using your knowledge from the Threads lab. You are going to use the client/server model in an upcoming assignment.

The lab includes a short introduction to sockets and the `java.net` package, which is going to help you write your own chat room application. You are going to implement a client that connects to a server, which maintains a list of all connected clients in the chat room. Your client is going to be able to communicate to all other connected clients by reading from and writing to a socket.

2 What is a Socket?

By definition a socket is one end of a two-way communication link between two programs running on the network. In other words, it is through sockets that clients access the network and transmit data.

On the server-side: Normally, the server runs on a specific computer and has a socket that is bound to a specific port number. The server just waits, listening to the socket, for a client to make a connection request on that port.

On the client-side: The client knows the hostname (this can be an IP address or a network ID such as `cslab1a`) of the machine on which the server is running and the port number on which the server is listening. To make a connection request, the client attempts to contact the server on the given hostname and port. The client also needs to give the server a way to contact it and so it binds to a local port number that it will use during this connection. This is usually assigned by the system.

If everything goes well, the server accepts the connection. Upon acceptance, the OS on the server machine creates a new socket with its remote endpoint set to the address and port of the client and the same local endpoint. A new socket is used so that the server can continue to listen to the original port for connection requests while tending to the needs of the newly-connected client.

Unfortunately, in the real world, sockets rarely go according to plan. The client, once it has initiated its connection, blocks all threads until a response has been received from the server. However, sometimes the server is busy or the request is lost over the network, and your program will be waiting forever. Just in case, it is good practice to use a timeout, which is the amount of time a client waits before giving up or trying to connect again. You can find more info on timeouts in the java docs. For more information on the uglier side of clients, take CS168.

3 Sockets in Java

In Java, socket classes are used to represent the connection between a client program and a server program. The `java.net` package provides two classes – `java.net.Socket` and `java.net.ServerSocket` – that implement the client side and the server side of the connection, respectively.

The `java.net.Socket` class is used by the client to establish a connection to the server. Connecting to the server can be accomplished by creating a `Socket` object. The socket at the client side just needs to know the host name (the name of the machine where server is running) and the port where the server is listening. Here is one way to create a socket:

```
Socket mySocket = new Socket("localhost", PortNumber);1
```

Establishing a server that monitors a particular port is done by creating an instance of the `java.net.ServerSocket` class. Here is one way to do that:

```
ServerSocket myServerSocket = new ServerSocket(PortNumber);
```

¹`localhost` is the hostname of the local machine.

The next step is to tell the newly created server to listen indefinitely and accept incoming requests. This is done by using the `accept()` method of the `ServerSocket` class. When a request comes, `accept()` returns a `Socket` object representing the connection. This is usually done in a loop.

Note: `accept` *blocks*. This means that execution will halt in this function call until it has created a value to return or receives a signal or interruption.

4 Reading from and Writing to a Socket

Communication between client and server is achieved by reading from and writing to a socket.

On the client side, communicating with the server is accomplished in two steps. First, the Input and Output stream corresponding to the `Socket` object are obtained. That can be done by using the `getInputStream()` and the `getOutputStream()` methods of the `Socket` class. In code it would be:

```
BufferedReader in = new BufferedReader(  
    new InputStreamReader(mySocket.getInputStream()));  
PrintWriter out = new PrintWriter(  
    mySocket.getOutputStream(), true);
```

Once these are set up, you can read and write using the corresponding streams. For example:

```
String line = in.readLine();  
out.println("Echo: " + line);
```

Note: `readLine` *blocks*. This means execution will halt in this function call until it has created a value to return or receives a signal or interruption.

On the server side, communicating with the client is done the same way. Since the communication has to continue until the client breaks the connection, the reading and writing is done within a loop. Once the client breaks the connection or stops sending the request, the `Socket` object representing the client has to be closed (similar to closing a file descriptor after you've opened it). This can be done by calling the `close()` method on the `Socket` object.

5 Write Your Own Client/Server Application

We provide you with the stencil code of a simple chatting application based on the client/server model. You should pull the stencil code from the repository. Your task is to read carefully the comments before each method and implement all of the methods. Methods that need to be filled out are marked with `TODO`'s. You are going to do this lab in two steps. First, you are going to implement the server, and once you feel you have it working, you are going to fill in the client. Your knowledge about threads is going to be useful in implementing the server since, in this model, a different thread is used to handle each client connection.

6 Part I: Implementing the Server Side

You are going to complete the server-side classes: `Server.java` and `IncomingClientHandler.java`. The server represents the chat room: it keeps track of all signed-on users, receives messages from them and sends the messages to all other users in the chat room.

The `Server.java` class has a `ServerSocket` which listens for connection requests from clients and accepts the connections. Once the connection is established, a new socket is associated with the client and a new `IncomingClientHandler` is created (check what `accept()` returns). The `ClientHandler` deals with receiving messages from the user and sending them to all users (including the sender).

7 Part II: Implementing the Client Side

For this part the class you have to complete is `Client.java`.

The `Client.java` class creates a new `Socket` and contains a thread for writing to and a thread for reading from the socket. Remember that in the constructor, you not only have to create the threads, but start them as well.

Once the client is connected to the server it has to be able to send messages to all other connected clients and receive messages from them. Sending a message is accomplished by first sending the message to the server, which maintains a list of all connected clients. The server then sends the message to the rest of the clients.

Note: The first message that should be sent to the server upon making a connection is the username passed in.

8 Running Your Chatroom

Unfortunately, NetBeans doesn't make it easy to run these client/server applications with the proper arguments. After using NetBeans to compile your code (use "Clean and Build"), you can use the provided *run-client* and *run-server* scripts (which can be found in this lab's directory in your cloned repository) to run the client and the server. The arguments you pass to *run-client* and *run-server* will be passed directly to your Java program.

If you need to debug, you will have to play with your NetBeans project's *Run Properties* (not fun - sorry!). For convenience, you can create multiple run configurations in NetBeans. These can be selected/activated from the drop-down box on the toolbar.

9 Testing

As always stress test your code and if possible make unit tests. Try running multiple clients and connecting to the same server, sending really long messages, blank messages etc.

Congratulations! You have your own chat room!

10 If you have time

You can try adding a prefix to all messages with the senders username. You can add sounds to play when messages are sent/received. Make an XML "buddy list" that is loaded and written to with perhaps a GUI. The possibilities are endless. Those commercial instant messaging clients guys have got nothing on you!